

Is Your Knowledge Ready for AI?

Before AI can answer from your organization's knowledge, that knowledge has to be findable, structured, traceable, owned, and checkable. Fifteen questions to check, before a model starts answering from it.

How to use it: pick **one real workflow or knowledge domain** — support answers, pricing rules, runbooks. Answer each question yes (Y), no (N), or not sure (?) for that domain. "Not sure" means nobody in the room knows — that counts as a finding.

1 · Findability

Can people and systems find the knowledge they need?

- ☐Y ☐N ☐? **Q1** Could a new teammate find the current answer to a routine question without asking anyone?
- ☐Y ☐N ☐? **Q2** Is there one agreed first place to look — or several plausible "sources of truth"?
- ☐Y ☐N ☐? **Q3** When an answer changes, does the old version get retired or clearly marked?

2 · Structure

Is the knowledge in a form AI systems can use?

- ☐Y ☐N ☐? **Q4** Could the core facts be laid out as a table, checklist, or set of rules without losing what matters?
- ☐Y ☐N ☐? **Q5** Is the same concept called the same name everywhere it appears?
- ☐Y ☐N ☐? **Q6** Are the decision rules written down as rules — or reconstructed from prose, screenshots, and threads?

3 · Provenance

Can you trace where knowledge came from?

- ☐Y ☐N ☐? **Q7** For a key document: can you tell who wrote it, when it was updated, and whether it's still current?
- ☐Y ☐N ☐? **Q8** Could you trace an AI answer back to the specific passage that supports it?
- ☐Y ☐N ☐? **Q9** When a significant decision gets made, is it recorded somewhere findable a year later?

4 · Ownership

Who maintains the knowledge?

- ☐Y ☐N ☐? **Q10** Can you name the person responsible for keeping each major source current?
- ☐Y ☐N ☐? **Q11** When two sources disagree, is there a defined way to resolve it?
- ☐Y ☐N ☐? **Q12** Has anything been updated this quarter for a reason other than something breaking?

5 · Verification

How would you know whether an AI answer is right?

- ☐Y ☐N ☐? **Q13** Do you have even ten questions with known-correct answers to test an AI against?
- ☐Y ☐N ☐? **Q14** Could your team agree — in writing — on which errors are annoying and which are unacceptable?
- ☐Y ☐N ☐? **Q15** If the AI gives a confident wrong answer in a high-stakes situation, who catches it, and can they override it?

How to read your answers — no score, just the pattern

Mostly yes: your knowledge may be ready for a narrow, well-scoped AI use case. Start small; keep Q13–Q15 answered in writing.

Mixed: the gaps will surface as "AI problems" if you build first. You likely need cleanup, modeling, or ownership decisions before a build.

Mostly no / not sure: you probably need a knowledge audit before relying on AI outputs. Building now means paying engineers to discover these gaps one incident at a time.

What to do next

1. Pick one knowledge domain — where AI help would matter most, not the tidiest one.
2. Identify the current source of truth. If there are several, choose one and say so.
3. List the contradictions and missing owners from Q10–Q12.
4. Write ten golden questions with expert-approved answers.
5. Decide what should not be automated yet, and write that down too.